

DATABASES

Lecture 7. Functions and Operators. Constraints.

Range Operators

Operator	Description	Example	Result
=	equal	<code>int4range(1,5) = '[1,4]':int4range</code>	t
<>	not equal	<code>numrange(1.1,2.2) <> numrange(1.1,2.3)</code>	t
<	less than	<code>int4range(1,10) < int4range(2,3)</code>	t
>	greater than	<code>int4range(1,10) > int4range(1,5)</code>	t
<=	less than or equal	<code>numrange(1.1,2.2) <= numrange(1.1,2.2)</code>	t
>=	greater than or equal	<code>numrange(1.1,2.2) >= numrange(1.1,2.0)</code>	t
@>	contains range	<code>int4range(2,4) @> int4range(2,3)</code>	t
@>	contains element	<code>'[2011-01-01,2011-03-01]':tsrange @> '2011-01-10':timestamp</code>	t
<@	range is contained by	<code>int4range(2,4) <@ int4range(1,7)</code>	t
<@	element is contained by	<code>42 <@ int4range(1,7)</code>	f
&&	overlap (have points in common)	<code>int8range(3,7) && int8range(4,12)</code>	t
<<	strictly left of	<code>int8range(1,10) << int8range(100,110)</code>	t
>>	strictly right of	<code>int8range(50,60) >> int8range(20,30)</code>	t
&<	does not extend to the right of	<code>int8range(1,20) &< int8range(18,20)</code>	t
&>	does not extend to the left of	<code>int8range(7,20) &> int8range(5,10)</code>	t
- -	is adjacent to	<code>numrange(1.1,2.2) - - numrange(2.2,3.3)</code>	t
+	union	<code>numrange(5,15) + numrange(10,20)</code>	[5,20)
*	intersection	<code>int8range(5,15) * int8range(10,20)</code>	[10,15)
-	difference	<code>int8range(5,15) - int8range(10,20)</code>	[5,10)

Range Functions

Function	Return Type	Description	Example	Result
<code>lower(anyrange)</code>	range's element type	lower bound of range	<code>lower(numrange(1.1,2.2))</code>	1.1
<code>upper(anyrange)</code>	range's element type	upper bound of range	<code>upper(numrange(1.1,2.2))</code>	2.2
<code>isempty(anyrange)</code>	boolean	is the range empty?	<code>isempty(numrange(1.1,2.2))</code>	false
<code>lower_inc(anyrange)</code>	boolean	is the lower bound inclusive?	<code>lower_inc(numrange(1.1,2.2))</code>	true
<code>upper_inc(anyrange)</code>	boolean	is the upper bound inclusive?	<code>upper_inc(numrange(1.1,2.2))</code>	false
<code>lower_inf(anyrange)</code>	boolean	is the lower bound infinite?	<code>lower_inf('(',')'::daterange)</code>	true
<code>upper_inf(anyrange)</code>	boolean	is the upper bound infinite?	<code>upper_inf('(',')'::daterange)</code>	true
<code>range_merge(anyrange, anyrange)</code>	anyrange	the smallest range which includes both of the given ranges	<code>range_merge('[1,2)'::int4range, '[3,4)'::int4range)</code>	[1,4)

Aggregate Functions

Function	Argument Type(s)	Return Type	Partial Mode	Description
<code>array_agg(expression)</code>	any non-array type	array of the argument type	No	input values, including nulls, concatenated into an array
<code>array_agg(expression)</code>	any array type	same as argument data type	No	input arrays concatenated into array of one higher dimension (inputs must all have same dimensionality, and cannot be empty or NULL)
<code>avg(expression)</code>	smallint, int, bigint, real, double precision, numeric, or interval	numeric for any integer-type argument, double precision for a floating-point argument, otherwise the same as the argument data type	Yes	the average (arithmetic mean) of all input values
<code>bit_and(expression)</code>	smallint, int, bigint, or bit	same as argument data type	Yes	the bitwise AND of all non-null input values, or null if none
<code>bit_or(expression)</code>	smallint, int, bigint, or bit	same as argument data type	Yes	the bitwise OR of all non-null input values, or null if none
<code>bool_and(expression)</code>	bool	bool	Yes	true if all input values are true, otherwise false
<code>bool_or(expression)</code>	bool	bool	Yes	true if at least one input value is true, otherwise false
<code>count(*)</code>		bigint	Yes	number of input rows
<code>count(expression)</code>	any	bigint	Yes	number of input rows for which the value of <i>expression</i> is not null
<code>every(expression)</code>	bool	bool	Yes	equivalent to <code>bool_and</code>
<code>json_agg(expression)</code>	any	json	No	aggregates values as a JSON array
<code>jsonb_agg(expression)</code>	any	jsonb	No	aggregates values as a JSON array
<code>json_object_agg(name, value)</code>	(any, any)	json	No	aggregates name/value pairs as a JSON object
<code>jsonb_object_agg(name, value)</code>	(any, any)	jsonb	No	aggregates name/value pairs as a JSON object

Aggregate Functions

<code>max(<i>expression</i>)</code>	any numeric, string, date/time, network, or enum type, or arrays of these types	same as argument type	Yes	maximum value of <i>expression</i> across all input values
<code>min(<i>expression</i>)</code>	any numeric, string, date/time, network, or enum type, or arrays of these types	same as argument type	Yes	minimum value of <i>expression</i> across all input values
<code>string_agg(<i>expression</i>, <i>delimiter</i>)</code>	(text, text) or (bytea, bytea)	same as argument types	No	input values concatenated into a string, separated by delimiter
<code>sum(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, numeric, interval, or money	bigint for smallint or int arguments, numeric for bigint arguments, otherwise the same as the argument data type	Yes	sum of <i>expression</i> across all input values
<code>xmlagg(<i>expression</i>)</code>	xml	xml	No	concatenation of XML values (see also Section 9.14.1.7)

`SELECT string_agg(name, ', ') FROM authors;`
Result: 'Tolstoy, Dostoevsky, Chekhov'

Date/time Functions

NOW() Get the current date and time	CURRENT_TIME Get the current time
CURRENT_DATE Get the current date	EXTRACT() Pull out date parts
AGE() Calculate time differences	INTERVAL Work with durations
DATE_TRUNC() Truncate to specific time units	TO_DATE() Convert strings to date
TO_CHAR() Format timestamps	CLOCK TIMESTAMP Real-time timestamps
TIMEOFDAY() Human-readable timestamp	TEMPOFDAY() Human-readable timestamp

CURRENT DATE & TIME

Examples

```
SELECT CURRENT_DATE; -- Result: 2025-10-14  
SELECT CURRENT_TIME; -- Result: 02:10:08.182677+06:00  
SELECT CURRENT_TIMESTAMP;  
-- Result: 2025-10-14 02:11:16.208125+06
```

- The **CURRENT_DATE** function will return the current date as a 'YYYY-MM-DD' format.
- **CURRENT_TIME** function will return the current time of day as a 'HH:MM:SS.GMT+TZ' format.
- The **CURRENT_TIMESTAMP** function will return the current date as a 'YYYY-MM-DD HH:MM:SS.GMT+TZ' format.

EXTRACT

EXTRACT function extracts parts from a date

Unit	Explanation
day	Day of the month (1 to 31)
decade	Year divided by 10
doy	Day of the year (1=first day of year, 365/366=last day of the year, depending if it is a leap year)
epoch	Number of seconds since '1970-01-01 00:00:00 UTC', if date value. Number of seconds in an interval, if interval value
hour	Hour (0 to 23)
minute	Minute (0 to 59)
month	Number for the month (1 to 12), if date value. Number of months (0 to 11), if interval value
second	Seconds (and fractional seconds)
year	Year as 4-digits

EXTRACT

EXTRACT function extracts parts from a date

Syntax

```
EXTRACT ( field FROM source )
```

Example

```
SELECT EXTRACT(DAY FROM TIMESTAMP '2025-10-14 21:30:45'); -- Result: 14
```

```
SELECT EXTRACT(MONTH FROM order_date) AS  
month,  
       COUNT(*) AS total_orders  
FROM orders  
GROUP BY 1  
ORDER BY 1;
```

Date_trunc()

`DATE_TRUNC` literally *cuts off* smaller time parts — for example, it rounds a timestamp to the start of a month, week, or day. It's a great way to group time data consistently.

Syntax

```
DATE_TRUNC ( 'field', source )
```

```
SELECT DATE_TRUNC('month', TIMESTAMP '2025-10-14  
15:27:34'); -- Result: 2025-10-01 00:00:00
```

Examples

```
SELECT DATE_TRUNC('month', order_date) AS month,  
COUNT(*) AS total_orders  
FROM orders  
GROUP BY 1  
ORDER BY 1;
```

EXTRACT VS. DATE_TRUNC

```
SELECT EXTRACT(MONTH FROM  
order_date) AS month,  
COUNT(*) AS total_orders  
FROM orders  
GROUP BY 1  
ORDER BY 1;
```

order_date	Result of EXTRACT(MONTH)
2024-01-10	1
2025-01-20	1
2025-02-05	2

```
SELECT DATE_TRUNC('month',  
order_date) AS month,  
COUNT(*) AS total_orders  
FROM orders  
GROUP BY 1  
ORDER BY 1;
```

order_date	Result of DATE_TRUNC('month')
2024-01-10	2024-01-01 00:00:00
2025-01-20	2025-01-01 00:00:00
2025-02-05	2025-02-01 00:00:00

AGE, INTERVAL CALCULATIONS

AGE function returns the number of years, months, and days between two dates.

Syntax

`AGE(timestamp1, timestamp2);`
If timestamp2 is NOT provided, current date will be used

`SELECT AGE('2025-10-14', '2000-02-01');` --25 years 8
mons 13 days

Examples

```
SELECT customer_name,  
       AGE(birth_date) AS age  
FROM customers  
ORDER BY age DESC;
```

ADDING AND SUBTRACTING TIME INTERVALS

Intervals can be expressed in many ways — days, hours, months, even combinations like '1 year 3 months 2 days'. PostgreSQL automatically adjusts the date.

Examples

```
SELECT CURRENT_DATE + INTERVAL '7 days' AS  
next_week;  
SELECT order_date + INTERVAL '30 days' AS  
delivery_date;  
SELECT CURRENT_DATE - INTERVAL '14 days' AS  
two_weeks_ago;
```

EXISTS

- The argument of EXISTS is an arbitrary SELECT statement, or subquery
- The subquery is evaluated to determine whether it returns any rows
- If it returns at least one row, the result of EXISTS is “true”
- If the subquery returns no rows, the result of EXISTS is “false”

EXISTS (subquery)

```
SELECT col1  
FROM tab1  
WHERE EXISTS (SELECT 1 FROM tab2 WHERE col2 = tab1.col2);
```


IN

- The right-hand side is a parenthesized subquery, which must return exactly one column
- The left-hand expression is evaluated and compared to each row of the subquery result
- The result of IN is “true” if any equal subquery row is found
- The result is “false” if no equal row is found

expression IN (subquery)

```
SELECT col1  
FROM tab1  
WHERE col1 IN (SELECT col2 FROM tab2);
```

NOT IN

- The right-hand side is a parenthesized subquery, which must return exactly one column
- The left-hand expression is evaluated and compared to each row of the subquery result
- The result of NOT IN is “true” if only unequal subquery rows are found
- The result is “false” if any equal row is found

`expression NOT IN (subquery)`

```
SELECT col1  
FROM tab1  
WHERE col1 NOT IN (SELECT col2 FROM tab2);
```

ANY/SOME

- The right-hand side is a parenthesized subquery, which must return exactly one column.
- The left-hand expression is evaluated and compared to each row of the subquery result using the given operator, which must yield a Boolean result.
- The result of ANY/SOME is “true” if any true result is obtained.
- The result is “false” if no true result is found

expression operator ANY (subquery)

expression operator SOME (subquery)

```
SELECT ProductName
```

```
FROM Products
```

```
WHERE ProductID = ANY (SELECT ProductID FROM OrderDetails WHERE Quantity = 10);
```

ALL

- The right-hand side is a parenthesized subquery, which must return exactly one column
- The left-hand expression is evaluated and compared to each row of the subquery result using the given operator, which must yield a Boolean result
- The result of ALL is “true” if all rows yield true
- The result is “false” if any false result is found

expression operator ALL (subquery)

```
SELECT ProductName
FROM Products
WHERE ProductID = ALL (SELECT ProductID FROM OrderDetails WHERE Quantity = 10);
```

CONSTRAINTS

- are used to define rules for columns in a database table. They ensure that no invalid data is entered into the database.
- gives as much control over the data in tables as you wish. If a user attempts to store data in a column that would violate a constraint, an error is raised.

CONSTRAINTS

- **NOT NULL:** a constraint that indicates that this column cannot be empty (NULL), and must have a value
- **PRIMARY KEY:** used to uniquely identify a row in the table. It is good practice to specify in each table the column that contains the primary key.
- **FOREIGN KEY:** Used to ensure referential integrity of the data.
- **UNIQUE:** Ensures that all values in a column are different.
- **CHECK:** Makes sure that all values in a column satisfy certain criteria.
- **DEFAULT:** you can set a default column value when no other value has been specified

CHECK

- A **CHECK** constraint is the most generic constraint type.
- It allows you to specify that the value in a certain column must satisfy a Boolean (truth-value) expression.

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CHECK (price > 0)  
);
```

CHECK

- Constraint definition comes after the data type, just like default value definitions.
- Default values and constraints can be listed in any order.

```
CREATE TABLE employees (  
  age INTEGER DEFAULT 25 CHECK (age >= 18));
```

```
CREATE TABLE employees (  
  age INTEGER CHECK (age >= 18) DEFAULT 25);
```

- A check constraint consists of the key word CHECK followed by an expression in parentheses.
- The check constraint expression should involve the column thus constrained, otherwise the constraint would not make too much sense.

```
age INTEGER CHECK (salary > 0)
```

CHECK

```
CREATE TABLE products_test (  
    product_no integer,  
    name text CHECK (char_length(name) > 3),  
    price numeric CHECK (price > 0)  
);
```

```
INSERT INTO products_test VALUES (1, 'abc', 100);
```

[23514] ERROR: new row for relation "products_test" violates check constraint "products_test_name_check"
Подробности: Failing row contains (1, abc, 100).

CHECK

```
CREATE TABLE products_test2 (  
    product_no integer,  
    name text CHECK (char_length(name) > 3) DEFAULT 'Hello',  
    price numeric DEFAULT 1 CHECK (price > 0)  
);
```

price numeric DEFAULT 0 CHECK (price > 0)

CHECK

- You can also give the constraint a separate name.
- This clarifies error messages and allows you to refer to the constraint when you need to change it.

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CONSTRAINT positive_price CHECK (price > 0)  
);
```

ERROR: new row for relation "products" violates check constraint "positive_price"

CHECK

- If you don't specify a constraint name in this way, the system chooses a name for you.
- Named constraints are especially helpful when you need to modify or drop them later.

`ALTER TABLE` products `DROP CONSTRAINT` positive_price;

- A check constraint can also refer to several columns.
- Say you store a regular price and a discounted price, and you want to ensure that

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CHECK (price > 0),  
    discounted_price numeric CHECK (discounted_price > 0),  
    CHECK (price > discounted_price)  
);
```


CHECK

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric,  
    CHECK (price > 0),  
    discounted_price numeric,  
    CHECK (discounted_price > 0),  
    CHECK (price > discounted_price)  
);
```

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric CHECK (price > 0),  
    discounted_price numeric,  
    CHECK (discounted_price > 0 AND price > discounted_price)  
);
```

CHECK

- Names can be assigned to table constraints in the same way as column constraints:

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric,  
    CHECK (price > 0),  
    discounted_price numeric,  
    CHECK (discounted_price > 0),  
    CONSTRAINT valid_discount CHECK (price > discounted_price)  
);
```

Not-Null

- A not-null constraint simply specifies that a column must not assume the null value.
- A not-null constraint is always written as a column constraint.
- A not-null constraint is functionally equivalent to creating a check constraint `CHECK (column_name IS NOT NULL)`
- But in PostgreSQL creating an explicit not-null constraint is more efficient.
- The drawback is that you cannot give explicit names to not-null constraints created this way.

Not-Null

- Column can have more than one constraint.
- Just write the constraints one after another:

```
CREATE TABLE products (  
    product_no integer NOT NULL,  
    name text NOT NULL,  
    price numeric NOT NULL CHECK (price > 0)  
);
```

- The NOT NULL constraint has an inverse: the NULL constraint.
- This simply selects the default behavior that the column might be null.

```
CREATE TABLE products (  
    product_no integer NULL,  
    name text NULL,  
    price numeric NULL  
);
```

Unique

- **UNIQUE** constraints ensure that the data contained in a column, or a group of columns, is unique among all the rows in the table.
- The syntax is:

```
CREATE TABLE products (  
    product_no integer UNIQUE,  
    name text,  
    price numeric  
);
```

```
CREATE TABLE products (  
    product_no integer,  
    name text,  
    price numeric,  
    UNIQUE (product_no)  
);
```

Unique

- To define a **UNIQUE** constraint for a group of columns, write it as a table constraint with the column names separated by commas:

```
CREATE TABLE example (  
    a integer,  
    b integer,  
    c integer,  
    UNIQUE (a, c)  
);
```

- The UNIQUE constraint allows NULL values.
- However, each NULL is treated as distinct, meaning multiple NULL values are allowed in a column with a UNIQUE constraint.
- Unlike the PRIMARY KEY constraint, which also enforces uniqueness, a table can have multiple UNIQUE constraints.

String Patterns

```
CREATE TABLE passengers (  
    passenger_id SERIAL PRIMARY KEY,  
    first_name VARCHAR(50) NOT NULL,  
    email VARCHAR(255),  
    CONSTRAINT check_email_format CHECK (  
        email ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'  
    )  
);
```

^[A-Za-z0-9._%+-]+: containing allowed characters.

@: symbol separating local part and domain.

[A-Za-z0-9.-]+: Domain name with allowed characters.

[A-Za-z]{2,}\$: Top-level domain with at least two letters.

```
INSERT INTO passengers (first_name, email) VALUES ('Alice', 'alice@gmail.com'); -- ✓
```

```
INSERT INTO passengers (first_name, email) VALUES ('Bob', 'not_an_email'); -- ✗
```

String Patterns

```
CREATE TABLE passengers_account(  
    passenger_id SERIAL PRIMARY KEY,  
    first_name VARCHAR(50) NOT NULL,  
    url varchar(255),  
    password VARCHAR(255),  
    CONSTRAINT check_password_complexity CHECK (  
        LENGTH(password) >= 8 AND  
        password ~ '[A-Z]' AND  
        password ~ '[a-z]' AND  
        password ~ '\d' AND  
        password ~ '[!@#$%^&*]'  
    )  
);
```

Add constraint

```
ALTER TABLE customers  
  ADD CONSTRAINT constr_name  
    CHECK(char_length(name) > 3)
```

```
ALTER TABLE customers  
  DROP CONSTRAINT constr_name;
```